



Department of Computer Science  
UNIVERSITY OF COLORADO **BOULDER**



# Machine Learning: Yoshinari Fujinuma

University of Colorado Boulder

LECTURE 17

Slides adapted from Chenhao Tan, Jordan Boyd-Graber

## Logistics

---

- Project proposal guideline will be released today

## Learning Objectives

---

- Recurrent Neural Network
- Convolutional Neural Network

## Outline

---

Recurrent Neural Networks (RNNs)

Convolutional Neural Network

## Motivation

---

- We used bag-of-words features for spam classification
- ...but what should we do when we want to distinguish “cats like dogs” vs. “dogs like cats”?

## Structured data

---

Sequential information

“My words fly up, my thoughts remain below: Words without thoughts never to heaven go.”

—Hamlet

## Structured data

---

### Sequential information

“My words fly up, my thoughts remain below: Words without thoughts never to heaven go.”

—Hamlet

- language
- activity history

## Structured data

---

### Sequential information

“My words fly up, my thoughts remain below: Words without thoughts never to heaven go.”

—Hamlet

- language
- activity history

Above sentence can be regarded as a sequence  $\mathbf{x} = (\mathbf{x}_1, \dots, \mathbf{x}_T)$ , starting from left ( $x_0 = \text{“My”}$ ) to right ( $x_T = \text{“go”}$ ).

e.g.,

$$\mathbf{x}_{\text{“My words fly up, my thoughts remain below”}} = (\mathbf{x}_{\text{My}}, \dots, \mathbf{x}_{\text{below}})$$



## Recurrent Neural Networks


---

Sharing parameters along a sequence

At time  $t$ , the hidden representation  $h_t$  is given as

$$h_t = f(x_t, h_{t-1})$$

i.e.,  $h_t$  is dependent on the hidden representation at the previous step  $t$ .



<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>

## Recurrent Neural Networks

---

In (vanilla) RNN, there are following training parameters

- parameter matrix on recurrent input  $W_{\text{rec}}$
- parameter matrix on input  $W_{\text{in}}$
- bias vector  $b_{\text{rec}}$  and  $b_{\text{in}}$

$$h_t = f(x_t, h_{t-1})$$

where


$$f(x_t, h_{t-1}) = g(W_{\text{in}}x_t + b_{\text{in}} + W_{\text{rec}}h_{t-1} + b_{\text{rec}})$$

and  $g$  is a non-linear activation function

## Recurrent Neural Networks

RNNs can be unrolled over time, good at capturing long-term dependencies.

E.g., open and closed brackets in C++ programming “if (condition) {...}”.



<https://colah.github.io/posts/2015-08-Understanding-LSTMs/>  
How to train this?

“Back propagation over time” and helps capture long-term dependencies

## Long short-term memory (LSTM)

---

Problem with (vanilla) RNN:

Vanishing gradient [Pascanu et al., 2013], hard to capture extremely long-term dependencies


## Long short-term memory (LSTM)

Problem with (vanilla) RNN:

Vanishing gradient [Pascanu et al., 2013], hard to capture extremely long-term dependencies

One hack: have “gates” to help remember/forget parameters

- $i_t \in \mathbb{R}^n$ : input gate at time  $t$
- $f_t \in \mathbb{R}^n$ : forget gate at time  $t$
- $o_t \in \mathbb{R}^n$ : output gate at time  $t$




## Long short-term memory (LSTM)

Problem with (vanilla) RNN:

Vanishing gradient [Pascanu et al., 2013], hard to capture extremely long-term dependencies

One hack: have “gates” to help remember/forget parameters


- $i_t \in \mathbb{R}^n$ : input gate at time  $t$
- $f_t \in \mathbb{R}^n$ : forget gate at time  $t$
- $o_t \in \mathbb{R}^n$ : output gate at time  $t$



## Long short-term memory (LSTM)

Intuitively, “gates” are vectors that controls the amount of inputs that can go through  
e.g., if we have a forget gate  $f$  at time  $t$  as  $f_t = (0, 0, 0, 0.1, 0)$ , then

$$f_t \odot (1, 2, 3, 4, 5) = (0, 0, 0, 0.4, 0)$$



$$f_t = \sigma(W_f \cdot [h_{t-1}, x_t] + b_f)$$

$$i_t = \sigma(W_i \cdot [h_{t-1}, x_t] + b_i)$$

$$o_t = \sigma(W_o \cdot [h_{t-1}, x_t] + b_o)$$

$$\tilde{C}_t = \tanh(W_C \cdot [h_{t-1}, x_t] + b_C)$$

$$C_t = f_t \odot C_{t-1} + i_t \odot \tilde{C}_t$$

$$h_t = o_t \odot \tanh(C_t)$$

## Pros and Cons of RNNs

---

- Pros
  - Can handle inputs with arbitrary length
- Cons
  - Slow due to going through one input after another
  - (and this is why Transformers [Vaswani et al., 2017] are popular these days)



## Outline

---

Recurrent Neural Networks (RNNs)

Convolutional Neural Network

## Structured data

---

### Spatial information




[https://www.reddit.com/r/aww/comments/6ip21a/before\\_and\\_after\\_she\\_was\\_told\\_she\\_was\\_a\\_good\\_girl/](https://www.reddit.com/r/aww/comments/6ip21a/before_and_after_she_was_told_she_was_a_good_girl/)

## Convolutional Layers

Sharing parameters across patches

input image  
or input feature map

output feature maps



$$a_{i'j'} = \sum_{i=1}^k \sum_{j=1}^k w_{ij} x_{ij}$$

- Number of filters
- Filter shape
- Stride size

<https://github.com/davidstutz/latex-resources/blob/master/tikz-convolutional-layer/convolutional-layer.tex>

## CNN

- Convolutional layer: extracting features, trainable parameters
- Pooling layer: downsampling, no parameters involved
- Full-connected layer: for optimizing on the objective function


E.g., convolutional layer with 4 filters, max pooling, and a fully-connected layer)

input image (10\*10)

4@6\*6

4@3\*3

5\*1



## Wrap up

---

### Neural Network architecture

- Recurrent Neural Networks
- Convolutional Neural Networks

## References

---

Razvan Pascanu, Tomas Mikolov, and Yoshua Bengio. On the difficulty of training recurrent neural networks. In *Proceedings of the 30th International Conference on Machine Learning*, pages 1310–1318, 2013. URL <http://proceedings.mlr.press/v28/pascanu13.html>.

Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N. Gomez, Lukasz Kaiser, and Illia Polosukhin. Attention is all you need. *CoRR*, abs/1706.03762, 2017. URL <http://arxiv.org/abs/1706.03762>.