Yoshinari Fujinuma

Office Address
Virtual

fujinumay@gmail.com
http://akkikiki.github.io
Twitter: @akkikiki

Education

University of Colorado Boulder, USA — Dec 2021
Ph.D. in Computer Science
Advisors
- Katharina Kann(2020-2021), Jordan Boyd-Graber(-2021), Michael J. Paul (-2019)

University of Tokyo, Japan — Sep 2014
M.S. in Computer Science; Advisor: Akiko Aizawa

International Christian University, Japan — Mar 2012
B.A. in Computer Science and Mathematics; Advisor: Grant Pogosyan

Professional Experience

Applied Scientist, Amazon.com, Inc., USA — Sep 2021 - Present
- AWS Bedrock: Worked on data curation and supervised fine-tuning for building the first-party LLM "Titan" using the Amazon Elastic Kubernetes Service (EKS) cluster.
- AWS Comprehend: Information extraction on document images/PDFs.

Part-Time Instructor, Univ. of Colorado Boulder, USA — Jan 2021 - May 2021
- CSCI 4622 Machine Learning (Github Repo)

Teaching Assistant, Univ. of Colorado Boulder, USA — Jan-Dec 2020, May-Jul 2021

Applied Scientist Intern, Amazon.com, Inc., USA — May 2020 - Aug 2020

Research Assistant, Univ. of Colorado Boulder, USA — Aug 2016 - Dec 2019

Applied Scientist Intern, Amazon.com, Inc., USA — May 2018 - Aug 2018

Software Dev Engineer, Amazon Japan K.K., Japan — Oct 2014 - Aug 2016

Software Engineer Intern, Amazon Japan K.K., Japan — Nov 2013 - Feb 2014

Part-time Engineer, Atilika, Japan — Aug- Nov 2013, Apr- Sep 2014

Software Engineer Intern, Cookpad, Japan — July 2013 (one month)

Publications

- Xiaoyu Liu, Huayang Li, Benjamin Hsu, Yoshinari Fujinuma, Maria Nadejde, Xing Niu, Ron Litman, Yair Kittenplon, Raghavendra Pappagari: "M^{3}T: A New Benchmark Dataset for Multi-Modal Document-Level Machine Translation", To appear at NAACL (short) 2024
- Yoshinari Fujinuma*, Siddharth Varia*, Nishant Sankaran, Srikar Appalaraju, Bonan Min, Yogarshi Vyas: "A Multi-Modal Multilingual Benchmark for Document Image Classification", EMNLP Findings (long) 2023
- Sharon Levy, Neha Anna John, Ling Liu, Yogarshi Vyas, Jie Ma, Yoshinari Fujinuma, Miguel Ballesteros, Vittorio Castelli, and Dan Roth: "Comparing Biases and the Impact of Multilingual Training across Multiple Languages", EMNLP (long) 2023

- Pietro Lesci, Yoshinari Fujinuma, Momchil Hardalov, Chao Shang, Lluis Marquez: "Diable: Efficient Dialogue State Tracking as Operations on Tables", ACL Findings (long) 2023
- Yoshinari Fujinuma, Jordan Boyd-Graber, Katharina Kann: "Match the Script, Adapt if Multilingual: Analyzing the Effect of Multilingual Pretraining on Cross-lingual Transferability", ACL (long) 2022
- Yoshinari Fujinuma, Masato Hagiwara: "Semi-Supervised Joint Estimation of Word and Document Readability", TextGraphs-15@NAACL (short) 2021
- Mozhi Zhang*, Yoshinari Fujinuma*, Michael J. Paul, Jordan Boyd-Graber: "Why Overfitting Isn't Always Bad: Retrofitting Cross-Lingual Word Embeddings to Dictionaries", ACL (short) 2020
- Mozhi Zhang, Yoshinari Fujinuma, Jordan Boyd-Graber: "Exploiting Cross-Lingual Subword Similarities in Low-Resource Document Classification", AAAI 2020
- Yoshinari Fujinuma, Jordan Boyd-Graber, Michael J. Paul: "A Resource-Free Evaluation Metric for Cross-Lingual Word Embeddings based on Graph Modularity", ACL (long) 2019
- Dasha Pruss, Yoshinari Fujinuma, Ashlynn R. Daughton, Michael J. Paul, Brad Arnot, Danielle Albers Szafir, Jordan Boyd-Graber: "Zika discourse in the Americas: A multilingual topic analysis of Twitter", PLOS ONE 2019
- Mozhi Zhang, Yoshinari Fujinuma, Jordan Boyd-Graber: "Exploiting Cross-Lingual Subword Similarities in Low-Resource Document Classification", Workshop on Deep Learning Approaches for Low-Resource Natural Language Processing, 2018
- Yoshinari Fujinuma, Alvin Grissom II: "Substring Frequency Features for Segmentation of Japanese Katakana Words with Unlabeled Corpora", IJCNLP (short) 2017
- Yoshinari Fujinuma, Hikaru Yokono, Pascual Martínez-Gómez, Akiko Aizawa: "Distant-supervised Language Model for Detecting Emotional Upsurge on Twitter", PACLIC (long) 2015

*denotes equal contribution

Selected Projects

Analysis on Multilingual Pretraining 2020 - 2021
- Investigating the effect of using different pretraining languages. (Paper)

Intrinsic Evaluation Measure for Cross-Lingual Embeddings 2017 - 2019
- Developed a graph-based intrinsic measure to evaluate the quality of cross-lingual word embeddings. (Paper)

Finite State Transducer (FST) for Kuromoji 2015
- Replaced a double-array trie to an FST for Kuromoji, a java-based Japanese tokenizer used in Lucene, Solr, and Elastic Search. (Github Repo)

Academic Service

Area Chair:
- 2023 ACL (Multilingualism & Cross-lingual NLP track)
- 2023 EMNLP (Multilingualism & Linguistic Diversity Track)
- 2023 LREC-COLING (Information Extraction, Knowledge Extraction, and Text Mining Track)

Program Committee/Reviewer:
- 2023 AAAI, ACL Rolling Review (Feb., Apr.), EACL, NAACL, W-NUT, NLP4HR
- 2022 AAAI, ACL Rolling Review (Jan., Mar., Jun., Oct., Dec.), ACL, EMNLP, NAACL SRW, W-NUT, CoNLL
- 2021 NAACL, ACL, EMNLP, ACL Rolling Review (Sept., Oct., Nov.), CoNLL, NAACL SRW, ACL SRW, W-NUT, Workshop on NLP for Indigenous Languages of the Americas (AmericasNLP)
- 2019 Workshop on Noisy User-Generated Text (W-NUT)

Secondary Reviewer: 2019 ACL, 2017 EMNLP, 2017 WWW

Student Volunteer: 2020 ACL

Academic Honors

- Outstanding Reviewer for CoNLL 2021
- Departmental Travel Grant from CU Boulder (300 USD) 2017
- Gruduate School Travel Grant from CU Boulder (500 USD) 2017
- Dean's Fellowship from CU Boulder (1/2 of the yearly tuition and stipend) 2016
- Gruduate School Travel Grant from Univ. of Tokyo (80,000 JPY) 2013
- Gödel Foundation Prize: Best Bachelor thesis in CS and Math (50,000 JPY) 2012
- Horie Takematsu and Koh Scholarship (2,000 USD) 2010
- Student Scholarship from ICU (1/3 of the yearly tuition) 2008-2012

Computer and Language Skills

Languages:	Proficient: Python; Intermediate: C++, Java
Libraries:	PyTorch, Deepspeed, MLFlow, Git, Vim, LaTeX
English:	TOEFL iBT 101 (2015)
Domain-specific:	machine learning, natural language processing